

Facilitating Large-Scale Graph Search Algorithms with Lock-Free Concurrent Pairing Heaps

Jeremy Mayeres Charles Newton Peter Tonner
jeremym@knights.ucf.edu newton@knights.ucf.edu ptonner@knights.ucf.edu

ABSTRACT

This paper introduces a lock-free version of a Pairing heap. Dijkstra’s algorithm is a search algorithm to solve the single-source shortest path problem. The efficiency of Dijkstra’s algorithm is asymptotically improved from $\mathcal{O}(|V|^2)$ to $\mathcal{O}(|E| + |V| \log(|V|))$ when an operation is available to decrease the recorded distance of a vertex to the target source in constant time (**decreaseKey**.) The performance of Dijkstra’s algorithm also improves when threads can also perform work concurrently (in particular, when **decreaseKey** calls occur concurrently.) However, current implementations of **decreaseKey** on popular backing data structures such as Pairing heaps and Fibonacci heaps severely limit concurrency. Lock-free techniques can improve the concurrency of search structures such as heaps. In this paper we introduce **decreaseKey** and **insert** operators for Pairing heaps that provide lock-free guarantees while still running in constant time. Techniques for parallelizing Dijkstra’s algorithm are additionally discussed.

Categories and Subject Descriptors

D.1.3 [Concurrent Programming]: Programming Techniques; E.1 [Lists, stacks, and queues]: Data Structures; E.1 [Trees]: Data Structures

General Terms

Performance, Algorithms, Parallel Algorithms

Keywords

Lock-free data structures, Pairing heap, Heap, Skiplist, Skip queue, Lock-free heap, Lock-free Pairing heap

1. INTRODUCTION

Developing an efficient concurrent lock-free self-balancing tree structure based on fine-grained synchronization primitives has so far remained an open problem [3] [11]. In this paper, we introduce various lock-free operations on Pairing

heaps [14], which are self-balancing tree structures that efficiently allow elements in the heap to be decreased.

As lock-free implementations of Pairing heaps are known to be difficult [17] and hitherto no attempts have been made to develop a linearizable concurrent Pairing heap based on fine-grained synchronization alone, we focus on only certain use cases. Specifically, our contribution is a lock-free concurrent Pairing heap capable of being a backing data structure for a parallel version of Dijkstra’s algorithm [8]. This is a reasonable simplification as the main motivation for Pairing heaps is efficiently running Dijkstra’s algorithm [14].

Dijkstra’s algorithm [8] is a search algorithm that solves the (single-source) shortest path problem for directed graphs with non-negative weights. It has wide applications in Internet routing (see, e.g., the shortest-path calculation in the OSPF routing protocol [4]) and other scheduling algorithms that depend on finding optimal paths.

Using a naïve data structure (such as a standard binary heap) results in Dijkstra’s algorithm having runtime of $\mathcal{O}(|V|^2)$, where $|V|$ is the number of vertices in the graph. Fredman and Tarjan have introduced [15] a heap variant, called the Fibonacci heap, where the **decreaseKey** operation takes $\mathcal{O}(1)$ time (amortized.) This allows Dijkstra’s algorithm to run in $\mathcal{O}(|E| + |V| \log(|V|))$ time, where $|E|$ is the number of edges in the graph. However, in practice, Fibonacci heaps have large constants that cause it to be slower than standard heap-backed priority queues on many practical graphs.

This deficiency led Fredman and Tarjan to develop the Pairing heap [14], which is simpler than Fibonacci heaps and has better performance in practice. Precise run-time bounds are, unfortunately, currently unknown. Pettie [39] has shown the **decreaseKey** operation takes between $\Omega(\log \log n)$ and $\mathcal{O}(2^{2\sqrt{\log \log n}})$ (amortized) time based on previous work by Fredman [13].

Current attempts to parallelize Dijkstra’s algorithm with lock-free data structures rely on **decreaseKey** operations that have an asymptotic performance worse than Pairing heaps. For example, to back Dijkstra’s algorithm with Shavit and Lotan’s concurrent priority queue [43] requires the user implement **decreaseKey** by deleting the target from the heap and reinserting it with a decreased value. Both of these operations run in $\mathcal{O}(\log n)$.

Input: A weighted graph $G = (V, E, W)$ and target node x

Output: The shortest distance from x to every vertex $v \in V$

```

1 begin
2   foreach  $v \in V$  do
3      $v.priority \leftarrow \infty$ ;
4      $v.previous \leftarrow \text{NULL}$ ;
5   end
6    $x.priority = 0$ ;
7    $PQ.insert(x)$ ;
8   while  $!PQ.empty$  do
9      $u \leftarrow deleteMin(PQ)$ ;
10    foreach  $v$  where  $(v, u) \in E$  do
11       $newDist \leftarrow u.priority + weight(v, u)$ ;
12      if  $newDist < v.priority$  then
13         $oldDist \leftarrow v.priority$ ;
14         $v.priority \leftarrow newDist$ ;
15         $v.previous \leftarrow u$ ;
16         $decreaseKey(oldDist - newDist, v, PQ)$ ;
17      end
18    end
19  end
20 end

```

Algorithm 1: Dijkstra’s algorithm. The main loop comprises a critical section which presents difficulties when parallelizing.

The main contributions of this paper are as follows. We introduce a lock-free variant of the Pairing heap that allows the **decreaseKey** and **insert** operations to be executed in parallel without the use of locks. Our lock-free **decreaseKey** operation has an asymptotic time complexity equal to Fredman and Tarjan’s [14] original **decreaseKey** implementation, modulo contention between threads.

In section 2 we discuss work in the literature related to concurrent constructions of heaps and priority queues. We give a review of Skiplists, both blocking and concurrent implementations, in section 3 and review blocking operations on Pairing heaps in section 4. A discussion of parallelization strategies for Dijkstra’s algorithm is given in section 5. We describe our lock-free Pairing heap implementation in section 6 and review its performance on real-world graph search problems in section 7. We give potential directions for later work on this problem in section 9.

2. RELATED WORK

Currently, the most common way to implement Dijkstra’s algorithm in a parallel manner is to partition the graph and apply Dijkstra’s algorithm on each subgraph [5]. In this section we outline some of the novel approaches to constructing concurrent priority queues previously found in the literature.

2.1 Concurrent Pairing Heaps

Hendler, Shavit, et al. [17] have applied their flat combining technique to construct a parallel Pairing heap. In their approach, each thread $0 \leq i < n$ places a description of its operation in the i^{th} index of an array. Threads race to acquire a global lock. The winning thread completes all n operations, combines the operations (e.g., several calls to decrease the value of the same node can be combined into one decrease

call), performs the operations serially, and writes the result of the i^{th} operation to the array. The losing threads wait to be notified that their operation has completed and return the result given by the winning thread. This approach, although concurrent, utilizes locks and, as a by-product, is vulnerable to undesirable side effects such as lock convoying and priority inversion [27].

2.2 Blocking Parallel Heaps

Nageshwara et al. [35] have implemented a priority queue with concurrent insert and delete operations using fine-grained locks. Their operations, although not lock-free, represent an attempt in the literature to create a concurrent heap model. Their model only scales to $\log(n)$ processors accessing a heap of n nodes. The binary heap implementation presented in their paper modifies the insert operation to work from top to bottom, rather than bottom up. This ensures that inserts can be run without deadlocks with concurrent delete operations, which also run top to bottom.

Driscoll et. al [9] have introduced a variant of Fibonacci heaps, called relaxed heaps, that allows for easier parallelization. However, their approach is complicated [10] and, in practice, performs worse than Fibonacci heaps [34].

2.3 Automatically Generated Lock-Free Heaps

In this section we consider how some approaches that automatically transform sequential data structures into lock-free, concurrent data structures could be applied to concurrent pairing heaps.

Herlihy [20] has introduced a universal method of automatically transforming sequential objects into concurrent, lock-free objects with very few assumptions about the behavior of the sequential object. In his method, all threads share a pointer to the sequential object and a version / timestamp of the object. When a concurrent operation is called, the object is copied into a new location in memory and operated upon with a corresponding sequential operation. The shared pointer is updated if the current timestamp matches the timestamp noted at the beginning of execution. Otherwise, the operation is repeated with a new copy of the object. This could lead to the construction of a lock-free pairing heap that copies itself into memory, performs the related sequential operation, and updates the state of the object to reflect these modifications (assuming no other thread has made changes.) However, as Herlihy notes, this automatic transformation is only suitable for objects which can be efficiently copied (i.e., “small objects.”) This would not be appropriate for our particular use-case since we are considering very large heaps.

Anderson and Moir [2] have improved on Herlihy’s approach by automatically keeping track of locations in memory a sequential operation reads from and writes to. Only modified blocks of memory are copied. Hence, this solves some issues with performance when only small portions of a data structure are modified. For example, consider updating the head of a linked-list backed queue. With Herlihy’s technique, the entire list would need to be copied, but the modification to the data structure only involves updating the head pointer. However, this method handles contention poorly if many

operations access the same blocks of memory (as is the case with operations defined on a Pairing heap.)

2.4 Concurrent Priority Queues and Heaps

Shavit and Lotan have constructed [43] a concurrent priority queue, called a SkipQueue, based on Pugh’s SkipList [41] data structure. Sundell and Tsigas [45] have improved on Shavit’s and Lotan’s method to create a fully lock-free concurrent priority queue.

Hunt et al. [24] have introduced an array-backed priority queue with very fine-grained locks. Their approach builds on previous work by Rao and Kumar [42] that showed contention on heaps can be greatly decreased by having the **deleteMin** and **insert** operations proceed in the same (top-down) directions. (If the operations proceed in opposite directions, this opens a possibility for deadlocks.) However, this increases the time complexity of an **insert** operation to $\mathcal{O}(\log n)$. The main contribution of Hunt et al.’s approach is that bottom-up **insert** operations can be preserved without introducing the possibility for deadlocks by 1) ordering the acquisition of locks, 2) applying tags / timestamps to avoid the ABA problem, and 3) avoiding overlap of search paths.

Israeli and Rappoport have developed [25] a wait-free concurrent heap. However, their implementation uses atomic operations that are not widely available in today’s hardware.

2.5 Concurrent Fibonacci and Skew Heaps

Jones [26] has constructed a concurrent Skew heap that uses fine-grained locks (as opposed to locking the entire heap structure) to allow parallel **insert** and **deleteMin** operations to execute without violating the heap invariant. In a Skew heap, each node has three pointers: two children pointers and one sibling pointer that points to the next node at the same depth. Jones defines a bubble as the set of all pointers a Skew heap operation will read or modify before completing its task. An operation must acquire all the locks in its bubble before performing any modifications. This granular lock structure allows disjoint heap operations to complete concurrently.

Huang and Weihl [23] have constructed a concurrent Fibonacci heap that allows threads to concurrently perform **insert**, **decreaseKey**, and **deleteMin** operations. The data structure is backed by a circular doubly-linked list. An issue with Huang and Weihl’s approach is that **deleteMin** is not linearizable: threads may access the minimum element of the Fibonacci heap out of order, (i.e., **deleteMin** does not necessarily return the minimum element in the heap at the point in time it is called), although the element is guaranteed to be reasonably small (“promising”) compared to the other elements in the heap. However, this presents problems with Dijkstra’s algorithm and other graph algorithms that require a stricter definition of **deleteMin**.

2.6 Parallel Implementations of Dijkstra’s Algorithm

In [37], Pearce et al. describe parallel implementations of many graph algorithms, including a single-source shortest path (SSSP) algorithm similar to Dijkstra’s algorithm. Their

approach to the SSSP is novel in that they eschew the need to synchronize the current node in the graph that has the least known distance to the target by, in some cases, making multiple visits to a node.

2.7 Other Highly Concurrent Attempts

Michael and Scott’s [32] lock-free FIFO queue allows enqueue and dequeue operations to proceed concurrently on a single link-list backed queue. Moir et al. have developed [33] a variant of the Michael-Scott queue that applies an elimination back-off technique to eliminate contention for the head / tail regions of the linked list. Previously, this technique was used by Shavit and Touitou [44] to create concurrent counters and Hendler et al. [18] to create concurrent LIFO stacks. Moir’s technique is a natural generalization of the elimination approach used in LIFO stacks: have the backed-off enqueue operations eliminate against a backed-off dequeue operation once all the elements prior to the back-off point have been dequeued.

3. SKIPLISTS

Skiplists were introduced by William Pugh [41] as a way to avoid expensive rebalancing in tree structures while still retaining desirable properties of balanced trees such as $\mathcal{O}(\log n)$ search. The basic idea behind Skiplists is to allow the tree structure to be controlled by a pseudorandom number generator (PRNG). Thus, a Skiplist only gives probabilistic guarantees that it will produce a balanced tree structure. However, pathologically bad structures that cause the data structure to degrade to a search complexity of $\mathcal{O}(n)$ are very unlikely.

A Skiplist is a set of linked lists with ordered elements. Each linked list is assigned a level $0 \leq i \leq M$ which describes the density of the linked list at that level. The first level (level 0) contains all the elements in the Skiplist. We adopt the notation that $L(i)$ denotes the list at level i . $L(0)$ (the bottom-most level) is guaranteed to contain all the elements inside the Skiplist. For each higher level $i > 0$, $L(i)$ is expected to contain $\frac{1}{2}|L(i-1)|$ elements. In Fig. 1, a Skiplist containing $L(0) = \{1, 3, 9, 11, 15\}$ is presented.




Figure 1: A Skiplist representation of $\{1, 3, 9, 11, 15\}$. The levels of each node, respectively, are 3, 0, 1, 2, and 0.

3.1 Core Operations

We now outline a few operations on Skiplists. In general, these are quite similar to their linked-list counterparts. In fact, each Skiplist operation can be decomposed into a sequence of linked-list operations that each affect one level of the Skiplist.

Suppose we want to search for element e . First, we look at the list at $L(M)$ and traverse it linearly until we encounter an element greater or equal to e . If we encounter e , then we stop. If we found an element greater than e , we repeat this procedure from this new starting point. If we only encountered the tail of the list, then we repeat the procedure starting from the first element in $L(M-1)$. This procedure repeats until we successfully locate e . If we encounter either an element greater than e or end of the list at $L(0)$ we conclude e is not contained in the Skiplist. See Figs. 2 and 3 for sample search runs on the Skiplist presented in Fig. 1.

To insert an element e into a Skiplist, we first perform a search for the element as previously described. However, we also keep track of additional information. The insertion point is determined by where we stop searching at $L(0)$. Throughout the search process, an array `previous[i]` is maintained that contains all the right-most pointers (at level i) to elements left of the insertion point. An array `next[i]` is given the previous values of each pointer in `previous[i]`. To insert e , we assign it a random level l , set all the right-most pointers in `previous[i]` for $0 \leq i \leq l$ to e and set the pointers emanating from e to `next[i]` for $0 \leq i \leq l$. Insertion has the visual effect of simply splicing e into the Skiplist (see Fig. 4.)

To delete an element e from a Skiplist, we first search for the element (as we did for insertion). If the element is not found, we are done. Otherwise, we should map each pointer in `previous[i]` to the value in `next[i]` for $0 \leq i \leq l$, where l is the level of e . See Fig. 4 for an illustration of deleting an element from a Skiplist.

3.2 Additional Operations

Pugh has defined [40] a number of useful operations such as merging Skiplists, splitting Skiplists, and concatenating two Skiplists together when all the keys in one Skiplist are less than or equal to the smallest key in a second Skiplist. However, an operation that decreases the value of an element in a Skiplist (besides simply deleting the old value and reinserting the decreased value) was not present and we were unable to find a description of this operation in the literature.




Figure 2: Skiplist traversal for finding the element 3 in the Skiplist presented in Fig. 1




Figure 3: Skiplist traversal for finding the element 15 in the Skiplist presented in Fig. 1




Figure 4: Illustration of insertion and removal. 5 has been inserted into the Skiplist. Red indicates pointers contained in `previous[i]`, and green indicates pointers contained in `next[i]`. 15 has also been removed from the Skiplist and updated pointers for this operation are in blue.

3.3 Lock-Free Skiplists

Shavit and Lotan's Skipqueue [43] is a modification of Pugh's Skiplist that gives its operations lock-free guarantees. We present a high-level technical summary of their work. Following with our analogy of Skiplist operations being composed of operations on the linked lists formed at each level of the Skiplist, we see that lock-free operations on Skiplists can be composed of operations on lock-free linked lists.

One subtle issue with Shavit and Lotan's modifications is that $L(i)$ is not necessarily a subset of all $L(k)$ (for $k < l$): a value's membership is determined by it being in the lowest level of the Skiplist. In particular, this means that even if we encounter a target value in a higher level of the Skiplist, we still must verify it is linked to the structure in $L(0)$.

The search, add, and delete operations of a lock-free Skiplist follow the same form as those we defined on Pugh's Skiplist. However, here we need to handle the possibility of threads concurrently modifying the Skiplist. To detect concurrent operations, we modify our linked list's node datatype from the 2-tuple $(\text{value}, \text{nextPtr}[])$ to the 3-tuple $(\text{value}, \langle \text{bool}, \text{nextPtr}[] \rangle)$, where `nextPtr` is the datatype used for pointers, and $\langle \text{bool}, \text{nextPtr}[] \rangle$ are written / read in a single atomic operation. In Java, this can be implemented with the `AtomicMarkableReference` class. In C/C++, this can be easily implemented by stealing a bit from a word-aligned pointer. The boolean value serves as a deletion bit that is set whenever a node is removed from the Skiplist but has yet to be physically removed from the linked list structure. This is important for avoiding certain interleavings of operations.

Consider two threads concurrently accessing a Skiplist. The first thread reads the deletion mark on an element a at level i , and is interrupted by a second thread which marks the element as deleted. When the first thread resumes, it will still assume that a is not marked and, if inserting a new element, could create an untraversable section of the Skiplist by building on a 's next pointer.

To search for an element in a lock-free Skiplist we take the same approach as before, but we additionally remove any marked nodes from the linked list structure and we always traverse down to $L(0)$ in order to verify membership. To add an element into the lock-free Skiplist we construct the `previous[i]` and `next[i]` sets as before and set the next pointers of the new node to `next[i]`. We also update the pointers in `previous[i]` to point to our new node (from level 0 to level l). If we are unsuccessful at having `previous[0]` point to our new node, we restart the whole procedure since either a new node was inserted and the ordering of $L(0)$ may be violated if we insert or the node at `previous[0]` was marked for removal. Note that the linearization point of an add occurs when the element is inserted into $L(0)$. Once `previous[0]` is updated, we attempt to update the remaining `previous[i]` to `next[i]` in a CAS loop. To handle concurrent changes during this process (e.g., `previous[i]` being marked), we recompute the `previous` and `next` arrays whenever an atomic update fails.

Removing an element from a lock-free Skiplist is analogous to adding an element. First, we search for the target element and compute our `previous[i]` and `next[i]` arrays. Next, we mark each `next` pointer of the target element in a top-down fashion within a CAS loop. When a CAS fails, we re-read the relevant `next[i]` as another thread could have modified it (e.g., a new node was linked to the node we're currently deleting.) When the `next` pointer has been marked in $L(0)$, we consider the element removed (this is our linearization point.) There is no need update any pointers on a removal since subsequent searches will perform this for us.

4. PAIRING HEAPS

Pairing heaps were developed by Fredman et al. [14] as a simpler alternative to Fibonacci heaps that, at least empirically, do not sacrifice performance. Pairing heaps are a popular backing data structure for Dijkstra's algorithm due to its offering of an efficient `decreaseKey` and its avoidance of complicated tree rebalances.

4.1 Operations on Pairing Heaps

We provide a sketch of the salient features of pairing heaps. The main operator defined on pairing heaps is `meld`, which accepts two heaps and uses a compare-and-link approach to join its inputs together. See Fig. 5 for an example of melding two heaps together. The roots are compared and the root with the largest value (assuming we have min-heaps) is added as the left-most child of the smaller root. Insertion of a node is performed by constructing a trivial heap containing only the relevant node and melding it the main heap.

The supported pairing heap operations are as follows:




Figure 5: Melding two min-heaps together (one with root 3 and the other with root 5) with the meld operator. The heap rooted at 5 becomes a subheap of the heap rooted at 3.

```
makeHeap(h)
findMin(h)
insert(x, h)
deleteMin(h)
meld(h, h')
decreaseKey( $\Delta$ , x, h)
delete(x, h)
```

where h, h' are pairing heaps, x is a target node (or, equivalently, a target value), and $\Delta > 0$ is the value by which the target node should be decremented.

`deleteMin` removes the minimum (root) element of a pairing heap H . Say once the root of the tree is removed we have n sub-trees H_1, H_2, \dots, H_n , each with their own individual roots. For each tuple in the set $\{(H_1, H_2), (H_3, H_4), \dots, (H_{n-1}, H_n)\}$, we meld the two heaps together to form $\{H'_1, \dots, H'_m\}$. It is worth mentioning that this pattern of grouping into tuples (pairs) is where pairing heaps derive their namesake. If n is odd, $H_n = H'_m$ is not melded. Finally, we perform a right-iteration (similar to foldr) of `meld` over these $m = \frac{n}{2}$ (or $m = \frac{n-1}{2} + 1$) heaps to construct a newly rebalanced pairing heap H' :

$$H' = \text{meld}(H'_1, \text{meld}(H'_2, \dots \text{meld}(H'_{m-1}, H'_m) \dots))$$

To make the ordering H_1, H_2, \dots, H_n unambiguous, we say that i^{th} child of H (i.e., the root of H_i) is the i^{th} most recently melded node to the root. Note here that `deleteMin` rebalances the tree structure of the pairing heap. See Fig. 6 for an illustration of calling `deleteMin` on a Pairing heap.

`decreaseKey` is implemented as follows. Suppose a node x has its value decreased and this new value violates the heap invariant (i.e., x is now smaller than its parent.) The subtree rooted at x is removed from the tree and melded to the parent tree as previously described.

`delete(x, h)` is implemented by delinking the sub-tree S rooted at x from h , calling `deleteMin` on S , and melding the resultant tree back to h .




Figure 6: Calling `deleteMin` on a Pairing heap (a). The root is removed (b) and the children are melded in pairs (c), and then the pairs themselves are melded in right-to-left order (d-e). Note that the resulting heap (e) is more balanced than (a).

5. PARALLELIZING DIJKSTRA’S ALGORITHM

We now consider some issues pertinent to parallelizing a heap-backed implementation of Dijkstra’s algorithm. If `deleteMin` does not remove and return the unvisited node in the graph with the smallest distance from the source node, Dijkstra’s algorithm is not guaranteed to solve the shortest-path problem. A naïve approach to running Dijkstra’s algorithm in parallel would be to allow calls to `deleteMin` and `decreaseKey` haphazardly interleave. However, the minimum of the heap will change on some call to `decreaseKey` if a new value of a node is less than the current value. Thus, some comparison is needed to see if calls to `decreaseKey` will produce a new minimum value in the heap. There are three basic approaches to solving this problem:

1. Optimistically continue removing minimum elements from the heap.
2. Pessimistically check if any of the new distances are less than the minimum of the heap, and continue if the minimum will remain unchanged. This has an advantage over the next method where the decision to continue can be made without any modifications to the heap.
3. Pessimistically assume that `decreaseKey` will always produce a new minimum and ensure that all calls to `decreaseKey` complete before the next iteration’s call to `deleteMin` executes.

Solution 1 requires complicated interactions between concurrent states of the heap or requires independent copies of the heap to be made and later merged. Since there are

two choices for every iteration of Dijkstra’s algorithm (either the heap’s minimum changes or does not change), a copying approach requires $\mathcal{O}(2^n)$ copies to be made, which is intractable for large graphs.

Solutions 2 and 3 are both practical. We opt for solution 3 since it is simpler, but for future work we may consider solution 2.

6. LOCK-FREE PAIRING HEAPS

In this section we develop lock-free `insert` and `decreaseKey` operators for Fredman et al.’s pairing heap. Our work relies on the observation that the specifics of Dijkstra’s algorithm allow us to make assumptions about how Pairing heap operations interleave. We first list some assumptions that we have made about overlapping calls to operations on the heap and give a high-level overview of our lock-free approach. In section 6.1 we give a technical overview of the difficulties encountered in developing a lock-free Pairing heap and give a high-level outline of our solutions to these issues. In section 6.2 we delve into the details of our `insert` and `decreaseKey` operators, in 6.3 we show our operations are linearizable and discuss the semantics of our lock-free pairing heap, and in 6.4 we sketch an informal proof regarding the correctness of our implementation.

In section 5 we analyzed the specific performance bottlenecks and other issues with running Dijkstra’s algorithm in parallel. In particular, we noted that the linearization guarantees provided by presently available concurrent heaps are insufficient for blindly executing concurrent calls to `deleteMin` and `decreaseKey`. Hence, we assume these calls will never interleave, and, further, that only concurrent calls to `decreaseKey` overlap.

6.1 Developmental Roadblocks

There are a number of issues that must be overcome when developing a lock-free Pairing heap. A design problem pertinent to all lock-free data structures is how to efficiently broadcast when one thread has made modifications that possibly interfere with concurrent work [19]. For a Pairing heap, we have identified four such coordination problems that arise from the need to maintain the heap root, the parent of a node in the heap, and the children of a heap node.

The first roadblock is that Pairing heap operations contend to make changes to the root of the heap when a node is produced (either a newly inserted node or an existing node that has sufficiently decreased) that is lower than all values currently in the heap. Nodes could be incorrectly deleted if `insert` and `decreaseKey` were to override the root of the heap without consideration of other changes. For example, if we concurrently `insert` two new values a, b (where $a < b$) into a heap that are both lower than all existing heap values, the intended state of the heap has a at the root, b as its child, and the remainder of the heap a child of either a or b . If a ’s insertion linearizes before b ’s insertion, the call to `insert` b must detect this and instead change its strategy to inserting b as a child of a . We have resolved this issue by using a CAS-loop that only commits changes to the heap root if no other changes have been made.

Another roadblock in designing a concurrent Pairing heap

is that calls to **decreaseKey** and **insert** that produce new roots compete to change the parent of the current root. Continuing along the same lines as our previous example, suppose that while inserting a and b , a linearizes first and hence has the old root as its child. If b concurrently tries to insert itself as the root, it will fail by the detection protocol established in the previous paragraph, but not before changing the root's parent from a to b . We resolve this problem by duplicating the old root of the heap and updating the parent field of the clone. This yields a useful property: if the new root fails to be CASed into place, no other node has access to the cloned previous root and it can safely be discarded.

However, we have introduced an additional problem: when the new root has been inserted, references to the old root must be atomically updated to the newly cloned node. In particular, this presents difficulties when maintaining references to nodes outside of the pairing heap structure. (For example, if an adjacency list of Pairing heap nodes is used in a graph, references must be updated.) We resolve this issue by first introducing an additional layer of indirection: all references to a Pairing heap node are made through a common reference object (pointer.) And, secondly, we ensure the process of updating the new location occurs atomically by the using a Descriptor-based design (see [12] and [7].) Descriptor objects combine the global state of an object and instructions for carrying out an operation. In our approach, Descriptors enclose the root node (the global state of the Pairing heap) and updating the memory location of the previous root (a pending operation.)

Our final coordination issue is that calls to **deleteMin** and **insert** also compete to insert nodes as a child of the current root. Say r is the root of the heap and we insert elements $a, b > r$. Both a and b will attempt to insert themselves as the left-most-child of r . We solve this issue by storing child nodes in a lock-free linked list (a Michael-Scott [32] queue.)

6.2 Lock-free insert and decreaseKey Operations

Our **insert** and **decreaseKey** operators utilize a Descriptor-based design: threads race to place Descriptors that describe operations to a Pairing heap. In general, our Descriptors are fine-grained and encapsulate changes to only one or two memory addresses.

Our **insert** operation first processes any pending operations. The insertion process is split into two cases. In the first case (lines 5-9 in Algorithm 2), the value of the root is less than or equal to the value of the node to be inserted, and the node is insert as the left-most child of the root.

In the second case (lines 10-19), the new node's value is less than the value of the root, and the root must be updated. To accomplish this update, we create a NewRoot-Descriptor that contains the new root, the old root, and the new memory address for the old root. We then try to CAS this descriptor into place. If CAS fails, we remove the root of the heap from the node we are inserting, process the current descriptor (if applicable), and retry.

Our **decreaseKey** operator is broken into four cases. In the first case (lines 5-21 of Algorithm 3), we are decreasing the

Input: A Pairing heap H and new node e
Output: e will be added to H

```

1 begin
2   while true do
3     // Case 1
4     Descriptor  $d \leftarrow \text{descriptor.get}()$ ;
5      $d.\text{execute}()$ ;
6     if  $e.\text{distance} \geq d.\text{root}.\text{distance}$  then
7        $e.\text{parent} \leftarrow d.\text{root}$ ;
8        $d.\text{root}.\text{subHeaps.add}(e)$ ;
9       break;
10    end
11    // Case 2
12    Node  $\text{rootClone} \leftarrow d.\text{root}.\text{clone}()$ ;
13     $e \leftarrow \text{meld}(\text{rootClone}, e)$ ;
14    Descriptor  $dNew \leftarrow \text{NewRootDescriptor}(e, d.\text{root}, \text{rootClone})$ ;
15    if CAS(descriptor,  $d$ ,  $dNew$ ) then
16      descriptor.get().execute();
17      break;
18    end
19    else
20       $e.\text{subHeaps.remove}(d.\text{root})$ ;
21      descriptor.get().execute();
22    end
23  end
24 end

```

Algorithm 2: A lock-free **insert** operator for Pairing heaps.

value of the current heap root. We first execute any pending operations, copy the current root, decrease the copy's value, and create a UpdateRootDescriptor that contains the new root and the memory location of the root of the heap. We then try to CAS the descriptor into place. If CAS succeeds, we are done. Otherwise, we retry after processing the descriptor if the target is still the root of the heap, or we move onto another case if the root has changed.

In the second case (lines 22-27), the new value of the node is still less than or equal to its parent's value. No structural changes to the heap are required in this case: we handle this case by simply changing the value of the node. One of our assumptions for **decreaseKey** is that calls to change a particular node do not interleave. A key benefit from this assumption is that this second case can proceed in a wait-free manner since no memory locations require an update.

Before executing the remaining two cases (lines 28-50), we remove the target node from its parent's list of subheaps as both cases involve moving the node to a new location in the heap. The final cases are largely identical to the cases of our **insert** operator as we are essentially inserting a new node into the heap. In the third case (lines 32-37), the node is larger than the current root and should be inserted as a child of the root. In the fourth case (lines 38-49), the new value of the target node is less than the value of the root and we insert the target node as the new root of the heap.

6.3 Linearization Points

Input: A Pairing heap H , a target node key , and a new weight w

Output: key 's value in the heap will be decreased to w .

```

1 begin
2   if !key.inHeap || key.distance ≤ w then
3     return;
4   end
  // Case 1
5   while key == descriptor.get().root do
6     Descriptor d ← descriptor.get();
7     d.execute();
8     if d.root != key then
9       break;
10    end
11    newRoot ← key.clone();
12    newRoot.distance ← w;
13    Descriptor dNew = UpdateRootDescriptor(d.root,
14    newRoot);
15    if CAS(descriptor, d, dNew) then
16      descriptor.get().execute();
17      return;
18    end
19    else
20      descriptor.get().execute();
21    end
22  end
  // Case 2
22  descriptor.get().execute();
23  key ← key.getPointer().getNode();
24  key.distance = w;
25  if key.parent.distance ≤ key.distance then
26    return;
27  end
28  key.parent.subHeaps.remove(key);
29  while true do
30    // Case 3
31    Descriptor d ← descriptor.get();
32    d.execute();
33    key ← key.getPointer().getNode();
34    if key.distance ≥ d.root.distance then
35      key.parent ← d.root;
36      d.root.subHeaps.add(key);
37      return;
38    end
39    // Case 4
40    key.parent ← null;
41    Node rootClone ← d.root.clone();
42    key ← meld(rootClone, key);
43    Descriptor dNew ← NewRootDescriptor(key,
44    d.root, rootClone);
45    if CAS(descriptor, d, dNew) then
46      descriptor.get().execute();
47      break;
48    end
49    else
50      key.subHeaps.remove(d.root);
51      descriptor.get().execute();
52    end
53  end
54 end

```

Algorithm 3: A lock-free decreaseKey operator for Pairing heaps.

```

1 if d.pending then
2   CAS (d.node, d.oldRoot, d.newRoot);
3   d.pending = false;
4 end

```

Algorithm 4: Execution of a NewRootDescriptor d .

```

1 if d.pending then
2   CAS (d.root, d.oldRoot, d.newRoot);
3   d.pending = false;
4 end

```

Algorithm 5: Execution of a UpdateRootDescriptor d .




Figure 7: Illustration of the concurrent pair heap structure. Pair heap nodes exist inside atomic references and hold atomic references to other heap nodes.

In this section we discuss the semantics of our lock-free Pairing heap and its linearization points. An operation is linearizable [22] if it atomically takes effect at some time t between the operation's invocation and its response. In our Pairing heap implementation, an operation linearizes when its changes to the heap take place.

Suppose we have n operations $\{\sigma_1, \dots, \sigma_n\}$, each either a call to **insert** or **decreaseKey**, concurrently executing. Let t_{desc}^i and t_{op}^i denote the time instants at which, respectively, σ_i atomically updates the Pairing heap's global descriptor object and when the memory writes encapsulated by σ_i 's descriptor object are executed. Note that a descriptor placed by σ_i may be executed by σ_i , some σ_j (where $j \neq i$), or completed by many σ_j 's each performing partial work on the descriptor. Also, let t_s^i denote the time where any changes to the list of subheaps of a particular node is modified by the operation σ_i .

We now turn to discussing the linearization points for an **insert** operation: an **insert** that involves placing the new node as a child of the current root linearizes at t_s^i when it is successfully placed into the list of subheaps (line 7 of Algorithm 2.) An **insert** that causes a root to be updated linearizes when its descriptor is placed at t_{desc}^i (line 13.) In both of these cases, the structure of the resultant heap was determined when the global descriptor was read at line 3.

As previously discussed, the final two cases for our **decreaseKey** operator are semantically identical to the two cases for **insert**. Hence, the linearization points for these two cases also align. By the same reasoning discussed in the previous paragraph, a **decreaseKey** operation falling into case three (lines 32-37 of Algorithm 3) linearizes at t_s^i (line 35), and an operation falling into case four linearizes at t_{desc}^i (line 42.) The heap's structure is determined by the particular descriptor read at line 30.

For the first case of **decreaseKey** (where we decrease the root of the heap), the operation linearizes when we place the descriptor into memory at t_{desc}^i (line 14) and the heap structure is determined at line 6. For the second case of **decreaseKey** (where the new value of the node is still greater than the parent and no structural changes are needed), the linearization point occurs when the value of the node is changed at line 24.

6.4 Correctness

In this section we discuss how our concurrent operations on Pairing heaps in all cases maintain the heap invariant and preserve the serial semantics of the **insert** and **decreaseKey** operations on Pairing heaps.

For serial Pairing heaps, an **insert** has two effects on the heap structure: either the new node becomes the new root of the Pairing heap (and the old root becomes its child) or the new node is added as the left-most child of the root of the Pairing heap. A **decreaseKey** has three effects: either the new node becomes (or is already) the heap root, is added as the left-most child of the first heap, or the target node does not change location at all.

We show that for all interleaving cases of **insert** / **de-**

creaseKey, the resulting modifications to the Pairing heap are consistent with some ordering of the same operations performed sequentially. Note that we continue with our assumption that concurrent calls to modify the same node in the heap do not occur. Note that both cases of **insert** are equivalent to the final two cases of **decreaseKey**, so we are free to omit discussion of **insert**.

Case 1. Suppose we are decreasing a root of the heap (lines 5-21 of Algorithm 3.) If another thread is decreasing an older root, then the descriptor has been modified and the CAS on line 14 will fail for that thread. When the descriptor is re-read, the current key will no longer be the root and the thread will be forced to complete their operation in a new case. If another thread is decreasing its heap to a value that is still less than its current parent (lines 23 - 27), then this does not involve any structural changes to the heap. The root modification can occur unaffected by this decrease. If the parent of the node being decreased is the root, since the reference to the parent node is made through a wrapper object (instead of a direct reference to the parent), the parent-child relationship will be preserved even though the root node is located at a new memory address (it is cloned, see line 11.)

The first interesting cases occur when we interleave a root update with a **decreaseKey** call that requires making changes near the root. Suppose a root update is interrupted by a second thread falling into the third case (the newly decreased node is made a child of the root.) However, no changes to the root are made in this case. The cloned root and the original root that exist in the root update (line 11) both share a reference to the same list of subheaps. Hence, the addition of the node is made immediately visible to both roots and successfully linearizes regardless of if the root is changed or not. A peculiar edge case here is when the root modification linearizes after the descriptor containing the old root has been read (line 14 completes after line 30.) The decreased node will be added on as a child of the old root, which is consistent with the Pairing heap state obtained when the child node decrease linearizes before the root modification linearizes. Otherwise, if the root modification linearizes before the descriptor has been read, the newly decreased child node will be added with the new root as its parent.

If the decreased node is less than the current root, there is a race to update the current root: either the decreased root will remain the root, or the newly decreased node will become the root. The linearization point for these cases is the CAS calls on lines 14 and 42, where the descriptor is updated. There are a few different cases to consider here. Let r and r' respectively denote the old and new values of the root (note $r' < r$) and say n is the new value of the node to be decreased. Since we are in lines 38-49, we additionally have $n < r$. Say that $r' < n < r$. If the new value of the root linearizes before the decreased node linearizes (if line 14 executes before line 42), then n will be the left-most child of r' . Otherwise, if the decreased node linearizes before the root, the CAS on line 14 will fail, and the root update will continue on to case 4. Next, suppose $n < r' < r$. If the newly decreased node n becomes the root before r' , then r' will be the left-child of n and the CAS on line 14 will fail. Since $n < r'$, r' will insert as a node in case 2. Note that

executing the descriptor on line 22 ensures that the old root's memory address is updated before we decrease its value on line 24. Otherwise, if the root's value changes to r' before n becomes the root, then the CAS on line 42 will fail and the descriptor to perform the memory address update from $r \rightarrow r'$ will be performed at line 31 (if it was not executed already by some competing thread.) Since $n < r'$, the node decrease will remain in case 4.

Case 2. Next we consider situations where operations (other than case 1 operations, which we have already considered) interleave with case 2 of `decreaseKey` (lines 22-27 of Algorithm 3.) Note that two competing case 2 operations on the same node cannot interleave by our assumption that no two calls to decrease the same node interleave. Also, if two case 2 operations interleave on non-connected parts of the heap, they clearly cannot interfere with each other. The case where we decrease a parent P with value p to $p' < p$ and a child C of P to $c' < c$ requires no structural changes. The linearization point for these interleavings when the value of the node is updated at line 24. The child node only proceeds to case 2 if $p \leq c'$ (if P 's change has not yet linearized) or of $p' \leq c'$ (if P 's change has linearized.) However, since $p' < p$ and $p \leq c'$ we have $p' < p \leq c'$ and C 's new value is a valid child of both p' and p . Hence, we can proceed without making any structural changes to the heap for these interleavings.

If a case 2 (lines 22-27) modification interleaves with a root modification (cases 3 and 4, lines 30-37 and lines 38-49), there are no possible concerns if we are modifying a node that is not P (adopting the same notation as before.) Further, if P is being modified in cases 3 or 4 while its child C is being modified in case 2, then the C is certainly still a valid child of P for the same reason, namely that P 's value will only decrease.

Case 3. We now consider interleavings with case three (lines 30-27 of Algorithm 3.) We have already covered two interleavings (cases 3 and 1 and cases 3 and 2.) If two concurrent calls to `decreaseKey` both fall into the third case, both will race to be the left-most child of the current root. The linearization point here is inside the `add` method of the Michael-Scott [32] linked-list based queue. The Pairing heap state (i.e., the ordering of the children) is consistent with the order in which these nodes reach this linearization point. If the fourth case executes concurrently with the third case, then there is a question about whether the node being decreased in the third case becomes a child of the old root or the new root. The descriptor read on line 30 contains the current root of the heap, and this determines which of the potential roots becomes the parent of the node being decreased. Note that even if the root changes between when the descriptor is read and when the node becomes a child of the root, the target node becomes a child of the node contained in the descriptor. Hence, the modifications to the Pairing heap state that result from the interleavings of cases 3 and 4 are always consistent with the moment in time the descriptor was read.

Case 4. Lastly, we consider the case where two newly decreased nodes are both less than the current value of the root and each contend to be the new root (case 4 - lines 22-27 of

Algorithm 3). Suppose the root has value r and the newly decreased nodes have values n and m such that $n, m < r$. Without a loss of generality, we assume that $n < m$ (the proof is the same for $m < n$, but with their respective roles reversed) The linearization point for case 4 is when the descriptor is placed into memory using CAS (line 42.) Suppose that the node with value m places its descriptor first. Then, the node with value n will fail to insert its descriptor and retry. Since $n < m$, the node with value n will again fall into case 4 and retry inserting its descriptor. When this succeeds, m will be its left-most child. Alternatively, suppose n inserts its descriptor first. Since $m > n$, m will retry in the third case (lines 30-37). Finally, we consider the case where $n = m$. Here, either node being the new root is a valid state, and the winner is determined by whichever node inserts its descriptor first on line 42. The remaining node will insert itself as a child of the winner (i.e., case 3.)

Note that since case 4 (lines 38-49) clones the current root of the heap, if the CAS call on line 42 fails, the old root is never modified and no other threads have access to this cloned node. This prevents a race condition where one thread overrides the parent field modification of another thread. Also note that our assumption that no two calls to `decreaseKey` change the same node ensures that the dangerous-looking changes to the list of subheaps on lines 40 and 47 are safe. In particular, we are certain to not, say, detach a node from a parent previously added by another call to `decreaseKey`.

7. EXPERIMENTAL RESULTS

In this section we compare the performance of lock-free Pairing heaps against Shavit and Lotan's Skipqueues [43], a priority queue based on a lock-free implementation of Pugh's [40] Skiplist. We ran a concurrent version of Dijkstra's algorithm (outlined in 5) that used Skipqueues and our lock-free Pairing heaps and compared the search completion time.

The graphs experimented against were obtained from the Stanford Large Network Dataset Collection [28]. We considered three datasets: a graph of the citations made between patents granted by the United States Patent and Trademark Office (USPTO) [30], a co-purchasing network graph from Amazon [29], and a graph representing links between UC Berkeley and Stanford's respective websites [31]. See Figure 7 for the number of vertices and edges contained in these large graph datasets.

To produce these results, we selected a node whose degree was equal to the mean degree of a vertex over the whole graph, rounded up, if needed. This ensured we did not select a node that biased our Pairing heap results due to its non-negligible start-up costs (see Section 8.) We searched on a varying number of threads (1, 2, 4, 8), and we report the average results obtained over 10 trials, respectively. These benchmarks were obtained using an Intel Xeon X5355 (2.66GHz), which has four physical cores and eight logical cores (via hyperthreading.)




Figure 12: Start-up state of a Pairing heap backing Dijkstra’s algorithm. The source node v has a distance of 0 and all other nodes $\{u_1, u_2, \dots, u_n\}$ are unvisited have currently have a value of ∞ .

Database Name	# Vertices	# Edges
Amazon Co-Purchase	262 111	1 234 877
Patent	3 774 768	16 518 948
BerkStanLinks	685 230	7 600 595

Figure 8: Sizes of graphs tested.

Results of the two competing strategies for searching the Amazon and patent datasets can be seen in Figure 9 and Figure 11. In general, the priority queue based on Skiplists out-performed the lock-free pairing heap. Significant anti-parallel effects were observed on the lock-free Pairing heap: aside from increasing the number of threads from 1 to 2 on the Amazon dataset, increasing the number of threads operating on the lock-free Pairing heap increased the graph search time. In contrast, Skiplists scaled well across many threads and had significantly less variance in their search times.

8. DISCUSSION

Our lock-free Pairing heap generally performed very poorly when compared with a Skiplist-backed heap. We now turn to discussing why this is the case and how this can direct future improvements on our initial algorithm design.

Consider the initial state of a Pairing heap constructed to back a search for the shortest distances from vertex v to all other vertices in the graph. The values in the heap will initially (see Algorithm 1) be $\text{dist}(v) = 0$ and $\text{dist}(u) = \infty$ for all other nodes $u \in \{u_1, u_2, \dots, u_n\}$. This implies that v is the root of the heap and all other nodes u are its children. As each u has an equal distance (∞), there is no real structure to the heap (see Figure 12.)

When we process the neighbors of v , we remove v from the heap as it is the minimum element. We replace it with some other unvisited element (say u_1 , see Figure 13.) When we decrease the distances of the neighbors of v , we delink them from their parents. In our current implementation, each node stores its subheaps in a Michael-Scott [32] queue. Unfortunately, this data structure has a linear search time. If $|v|$ denotes the degree of v , we have a total complexity of $\mathcal{O}(|v|^2)$, which is not ideal. In a Pairing heap, this work is amortized over the life of the data structure: early calls to `decreaseKey` occur in, at minimum, $\mathcal{O}(|v|)$ time, but later calls happen in near constant time. However, if $|v|$ is quite large, then an unreasonable amount of time is spent performing the first `decreaseKey`.

We also observed a decrease in performance when the number of threads working on `decreaseKey` operations were in-




Figure 13: State of a Pairing heap after running `deleteMin()` on the heap in Fig. 12.

creased. In blocking Pairing heaps, nodes are inserted (a similar procedure is used for `decreaseKey`) into the heap by either replacing the root of the heap or inserting itself as a subheap to the root. A parallel variant Pairing heap variant that preserves the general form of insertion operations will place un-due contention near the root of the data structure. Always inserting at the root appears to be a significant bottleneck. Skipqueues, on the other hand, can insert elements in, e.g., the middle of the list in a less-contended way and are thus more parallel.

9. FUTURE WORK

In this section we outline potential future directions for further research in developing efficient lock-free pairing heaps.

9.1 Additional Operators

We briefly outline a sketch of a parallelized `deleteMin` operator on Pairing heaps. `deleteMin` follows what is essentially a Map-Reduce pattern: pairs of heaps are constructed and then merged together to form a resultant heap. The first step has no contention between pairs and is trivial to parallelize. Our `meld` operator can be used to merge heaps in parallel.

However, once a `deleteMin` process has started, the new heap root cannot be determined until the call has finished. To accomplish `deleteMin` in a lock-free manner, a Descriptor object could be placed into memory that forces calls to other heap operations to first help complete the `deleteMin` before proceeding to their own operations.

9.2 Relieving Contention

Our current implementations of `insert` and `decreaseKey` place heavy contention on two memory addresses: the location containing the root of the heap and the left child of the current root. The first contention arises from threads competing to change the root of the heap, and the second from threads placing nodes under the root. To prevent these contentious situations, an elimination strategy (similar to [18]) could be employed. As for two different node values a, b either $a \leq b$ or $b < a$, two contending nodes could be combined using the `meld` operator described in section 4. The greater of the two nodes would be added as a child of the lesser, and only the lesser would need to contend for insertion near (or at) the root.

9.3 Avoiding Large Start-up Costs

A problem discussed in Section 8 is that the initial state of the Pairing heap lends itself to a $\mathcal{O}(|v'|)$ start-up time, where $|v'|$ denotes the number of edges connected to the source node. There are two clear ways to resolve this issue. First, we could artificially introduce structure onto the initial heap state (see Figure 14) to reduce the search space for

Dijkstra's Algorithm Performance (s) with Skiplists and Lock-free Pairing Heaps. Amazon Dataset.




Figure 9: Completion time of Dijkstra's algorithm on the Amazon co-purchasing network [29]. Error bars represent standard deviation over 10 trials.

Dijkstra's Algorithm Performance (s) with Skiplists and Lock-free Pairing Heaps. Patent Dataset.




Figure 10: Completion time of Dijkstra's algorithm on the USPTO database [30]. Error bars represent standard deviation over 10 trials.

Dijkstra's Algorithm Performance (s) with Skiplists and Lock-free Pairing Heaps. Stanford Berkeley Link Dataset




Figure 11: Completion time of Dijkstra's algorithm on the Berkeley / Stanford link dataset [31]. Error bars represent standard deviation over 10 trials.




Figure 14: A more structured heap (see Figure 12.) Note that the heap invariant is not violated by a parent and child having equal weights.

a particular node. A second solution would be to replace the data structure that represents the list of subheaps. For example, we could replace this with a Skiplist (which offers $\mathcal{O}(\log n)$ search) or use a hash table to insert search fingers [16] into our linked list.

9.4 Reliable Benchmarking

At the moment, our implementation of a lock-free Pairing heap is written in Java 1.7. The Java virtual machine (JVM) performs various tasks in the background [38] such as freeing unused memory, runtime code optimization [36], and heap management [1]. These tasks are unpredictable and may have introduced additional bias and variance in our experimental results. A solution to these problems is to reimplement our prototype Pairing heap in a language such as C++ that gives explicit control over memory management and performs all code optimizations at compile time.

9.5 ABA Problem

We have not considered the ABA problem in all of our design. Although the descriptor object is guarded by Doug Lea's Java implementation a timestamped atomic reference (`AtomicStampedReference`) [21], the ABA problem is not handled when the descriptor object is executed. To remedy this, we should apply coping techniques for accounting for the ABA problem in descriptor-based designs [6] [7].

10. CONCLUSION

We presented an implementation of a lock-free concurrent Pairing priority heap that allows multiple threads to interleave insertions and heap modifications. The concurrent Pairing heap is functional enough to back a parallelized version of Dijkstra's algorithm. Although the lock-free Pairing heap is currently not competitive with state-of-the-art concurrent priority heap structures such as Skipqueues, we have outlined future directions for improvement. We additionally have made progress in developing methods for constructing highly concurrent self-balancing tree structures with lock-free progress guarantees.

11. REFERENCES

- [1] E. Altman, M. Arnold, R. Bordawekar, R. M. Delmonico, N. Mitchell, and P. F. Sweeney. Observations on tuning a java enterprise application for performance and scalability. *IBM J. Res. Dev.*, 54(5):494–505, Sept. 2010.
- [2] J. Anderson and M. Moir. Universal constructions for large objects. *Parallel and Distributed Systems, IEEE Transactions on*, 10(12):1317–1332, dec 1999.
- [3] N. G. Bronson, J. Casper, H. Chafi, and K. Olukotun. A practical concurrent binary search tree. *SIGPLAN*

Not., 45(5):257–268, Jan. 2010.

- [4] R. Coltun, D. Ferguson, J. Moy, and A. Lindem. OSPF for IPv6. RFC 5340 (Proposed Standard), July 2008.
- [5] A. Crauser, K. Mehlhorn, U. Meyer, and P. Sanders. A parallelization of dijkstra’s shortest path algorithm. In L. Brim, J. Gruska, and J. Zlatuška, editors, *Mathematical Foundations of Computer Science 1998*, volume 1450 of *Lecture Notes in Computer Science*, pages 722–731. Springer Berlin / Heidelberg, 1998. 10.1007/BFb0055823.
- [6] D. Dechev. The ABA problem in multicore data structures with collaborating operations. In *CollaborateCom*, pages 158–167, 2011.
- [7] D. Dechev, P. Pirkelbauer, and B. Stroustrup. Understanding and effectively preventing the ABA problem in descriptor-based lock-free designs. In *Proceedings of the 2010 13th IEEE International Symposium on Object/Component/Service-Oriented Real-Time Distributed Computing, ISORC ’10*, pages 185–192, Washington, DC, USA, 2010. IEEE Computer Society.
- [8] E. W. Dijkstra. A note on two problems in connexion with graphs. *Numerische Mathematik*, 1:269–271, 1959. 10.1007/BF01386390.
- [9] J. R. Driscoll, H. N. Gabow, R. Shrairman, and R. E. Tarjan. Relaxed heaps: an alternative to fibonacci heaps with applications to parallel computation. *Commun. ACM*, 31:1343–1354, November 1988.
- [10] A. Elmasry. The violation heap: A relaxed fibonacci-like heap. In M. Thai and S. Sahni, editors, *Computing and Combinatorics*, volume 6196 of *Lecture Notes in Computer Science*, pages 479–488. Springer Berlin / Heidelberg, 2010. 10.1007/978-3-642-14031-0_51.
- [11] K. Fraser. *Practical lock freedom*. PhD thesis, Cambridge University Computer Laboratory, 2003. Also available as Technical Report UCAM-CL-TR-579.
- [12] K. Fraser and T. Harris. Concurrent programming without locks. *ACM Trans. Comput. Syst.*, 25(2), May 2007.
- [13] M. L. Fredman. On the efficiency of pairing heaps and related data structures. *J. ACM*, 46:473–501, July 1999.
- [14] M. L. Fredman, R. Sedgewick, D. D. Sleator, and R. E. Tarjan. The pairing heap: a new form of self-adjusting heap. *Algorithmica*, 1:111–129, January 1986.
- [15] M. L. Fredman and R. E. Tarjan. Fibonacci heaps and their uses in improved network optimization algorithms. *J. ACM*, 34:596–615, July 1987.
- [16] E. N. Hanson and T. Johnson. The interval skip list: A data structure for finding all intervals that overlap a point. In *In Proc. of the 2nd Workshop on Algorithms and Data Structures*, pages 153–164. Springer, 1992.
- [17] D. Hendler, I. Incze, N. Shavit, and M. Tzafrir. Flat combining and the synchronization-parallelism tradeoff. In *Proceedings of the 22nd ACM symposium on Parallelism in algorithms and architectures, SPAA ’10*, pages 355–364, New York, NY, USA, 2010. ACM.
- [18] D. Hendler, N. Shavit, and L. Yerushalmi. A scalable lock-free stack algorithm. In *Proceedings of the sixteenth annual ACM symposium on Parallelism in algorithms and architectures, SPAA ’04*, pages 206–215, New York, NY, USA, 2004. ACM.
- [19] M. Herlihy. Wait-free synchronization. *ACM Trans. Program. Lang. Syst.*, 13(1):124–149, Jan. 1991.
- [20] M. Herlihy. A methodology for implementing highly concurrent data objects. *ACM Trans. Program. Lang. Syst.*, 15:745–770, November 1993.
- [21] M. Herlihy and N. Shavit. *The Art of Multiprocessor Programming*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2008.
- [22] M. P. Herlihy and J. M. Wing. Linearizability: a correctness condition for concurrent objects. *ACM Trans. Program. Lang. Syst.*, 12(3):463–492, July 1990.
- [23] Q. Huang and W. Weihl. An evaluation of concurrent priority queue algorithms. In *Parallel and Distributed Processing, 1991. Proceedings of the Third IEEE Symposium on*, pages 518–525, dec 1991.
- [24] G. Hunt, M. M. Michael, S. Parthasarathy, and M. L. Scott. An efficient algorithm for concurrent priority queue heaps. *Inf. Proc. Letters*, 60:151–157, 1996.
- [25] A. Israeli and L. Rappoport. Efficient wait-free implementation of a concurrent priority queue. In A. Schiper, editor, *Distributed Algorithms*, volume 725 of *Lecture Notes in Computer Science*, pages 1–17. Springer Berlin / Heidelberg, 1993. 10.1007/3-540-57271-6_23.
- [26] D. W. Jones. Concurrent operations on priority queues. *Commun. ACM*, 32:132–137, January 1989.
- [27] S. Kleiman, D. Shah, and B. Smaalders. *Programming with threads*. Sun Soft Press, 1996.
- [28] J. Leskovec. Stanford large network dataset collection, 2012.
- [29] J. Leskovec, L. A. Adamic, and B. A. Huberman. The dynamics of viral marketing. *ACM Trans. Web*, 1(1), May 2007.
- [30] J. Leskovec, J. Kleinberg, and C. Faloutsos. Graphs over time: densification laws, shrinking diameters and possible explanations. In *Proceedings of the eleventh ACM SIGKDD international conference on Knowledge discovery in data mining, KDD ’05*, pages 177–187, New York, NY, USA, 2005. ACM.
- [31] J. Leskovec, K. J. Lang, A. Dasgupta, and M. W. Mahoney. Community structure in large networks: Natural cluster sizes and the absence of large well-defined clusters. *Internet Mathematics*, 6(1):29–123, 2009.
- [32] M. M. Michael and M. L. Scott. Simple, fast, and practical non-blocking and blocking concurrent queue algorithms. In *Proceedings of the fifteenth annual ACM symposium on Principles of distributed computing, PODC ’96*, pages 267–275, New York, NY, USA, 1996. ACM.
- [33] M. Moir, D. Nussbaum, O. Shalev, and N. Shavit. Using elimination to implement scalable and lock-free fifo queues. In *Proceedings of the seventeenth annual ACM symposium on Parallelism in algorithms and architectures, SPAA ’05*, pages 253–262, New York, NY, USA, 2005. ACM.
- [34] B. M. E. Moret and H. D. Shapiro. An empirical assessment of algorithms for constructing a minimum

spanning tree, 1994.

- [35] R. Nageshwara and V. Kumar. Concurrent access of priority queues. *Computers, IEEE Transactions on*, 37(12):1657–1665, Dec 1988.
- [36] K. Ogata, D. Mikurube, K. Kawachiya, S. Trent, and T. Onodera. A study of java’s non-java memory. *SIGPLAN Not.*, 45(10):191–204, Oct. 2010.
- [37] R. Pearce, M. Gokhale, and N. M. Amato. Multithreaded asynchronous graph traversal for in-memory and semi-external memory. In *Proceedings of the 2010 ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis*, SC ’10, pages 1–11, Washington, DC, USA, 2010. IEEE Computer Society.
- [38] T. Peierls, B. Goetz, J. Bloch, J. Bowbeer, D. Lea, and D. Holmes. *Java Concurrency in Practice*. Addison-Wesley Professional, 2005.
- [39] S. Pettie. Towards a final analysis of pairing heaps. In *Foundations of Computer Science, 2005. FOCS 2005. 46th Annual IEEE Symposium on*, pages 174 – 183, oct. 2005.
- [40] W. Pugh. A skip list cookbook. Technical report, University of Maryland, 1990.
- [41] W. Pugh. Skip lists: a probabilistic alternative to balanced trees. *Commun. ACM*, 33:668–676, June 1990.
- [42] V. N. Rao and V. Kumar. Concurrent access of priority queues. *IEEE Trans. Comput.*, 37:1657–1665, December 1988.
- [43] N. Shavit and I. Lotan. Skiplist-based concurrent priority queues. In *Parallel and Distributed Processing Symposium, 2000. IPDPS 2000. Proceedings. 14th International*, pages 263–268, 2000.
- [44] N. Shavit and D. Touitou. Elimination trees and the construction of pools and stacks. *Theory of Computing Systems*, 30:645–670, 1997. 10.1007/s002240000072.
- [45] H. Sundell and P. Tsigas. Fast and lock-free concurrent priority queues for multi-thread systems. *J. Parallel Distrib. Comput.*, 65:609–627, May 2005.