Facilitating Large-Scale Graph Searches with Lock-Free Pairing Heaps

Jeremy Mayeres, Charles Newton, Peter Tonner

12 March 2012

Project Overview

- We are constructing a lock-free version of a Pairing heap (a self-balancing heap.)
- Pairing heaps have an efficient decreaseKey implementation (near-constant performance) that allow you to decrease the value in a heap without reinserting it.
- This improves the asymptotic performance of certain algorithms (e.g., Dijkstra's algorithm.)
- We are comparing our heap against Skipqueues (a priority queue backed with a Skiplist.)

Dijkstra's Algorithm

Single-Source Shortest Path Problem

In a weighted graph $G = (V, E)$, find the shortest path from a target vertex $v \in V$ to all other vertices in the graph.

- Push every node in a graph into a priority queue (PQ.) In the PQ, each node is weighted by current information we have about their distance.
- Initially, all nodes have a distance of ∞, except the target node, which has a distance of 0.
- Dynamic programming approach: Inductively build up our shortest routes.
  - Pop off the node on the PQ with the smallest distance. (The shortest path from the source to this node is finished.)
  - Update the weights in the PQ with the distances emanating from the popped node.

Skiplists

- Skiplists are constructed from a hierarchy of linked lists with the *Skiplist property*: the set of elements contained in level i is a subset of all the levels below it.
- Links allow you to do a binary search and "jump" around a list.
- The height of an element is randomly sampled from a power-law distribution: the probability of a node having a height of $i \geq 0$ is 2^{-i}.
- Time complexity of operations are probabilistically the same as for a binary search tree, but $\mathcal{O}(n)$ in the worst case.
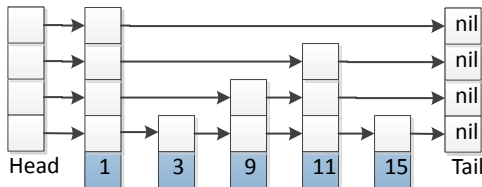
Skiplists (Search)

- For each level, move right until you run into a node greater than your target. Then, from this point, move right on the next lowest level, and repeat.
- Likely runs in $\mathcal{O}(\log n)$ time.
- Head / tail nodes have values of $-\infty$ and ∞, respectively.

Next and Previous Windows

- A useful abstraction to make is the set of all pointers related to a given node. (Aggregate all the pointers related to one node.)
- $prev[i] \rightarrow$ the node in level i pointing to the target node.
- $next[i] \rightarrow$ the node the target node points to at level i.
- Use the search process to construct these sets.
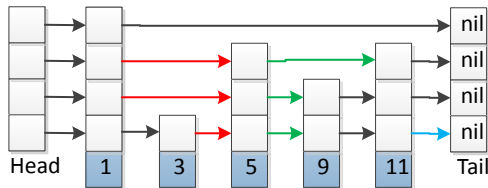
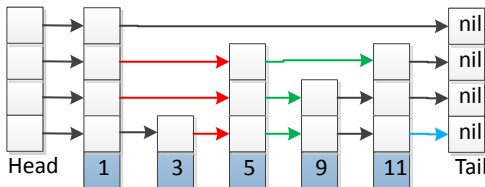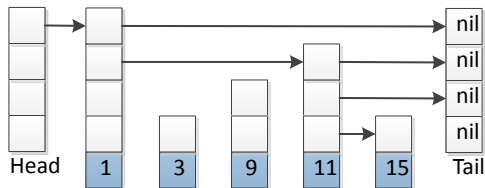## Skiplists (Insertion)

- Insertion: Find the location the target should be at. Set all of the node's next pointers to $next[i]$ and set the next pointers at each $prev[i]$ to the new node.
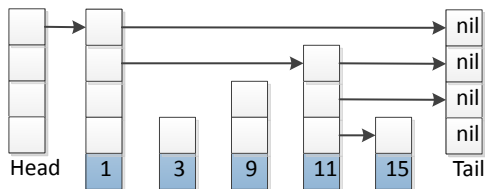
Skiplists (Deletion)

- Deletion: Set all of the next pointers in prev[i] equal to next[i].

Lock-free Skiplists

- Instead of linked lists at each level, we use lock-free linked lists.
- We lose the Skiplist property. (Levels aren't necessarily subsets of each other.) In particular, this means we need to always verify a node is in the lowest level.

Lock-free Skiplists

- Instead of linked lists at each level, we use lock-free linked lists.
- Linearization point: membership is defined at the lowest level of the Skiplist.
- We lose the Skiplist property. (Levels aren't necessarily subsets of each other.) In particular, this means we need to always verify a node is in the lowest level.

Lock-Free Skiplists (Insertion)

- Construct the prev[i] and next[i] sets.
- Set all of the node's next pointers to next[i].
- If we can CAS the node into the bottom level, continue. Otherwise, something changed, and restart (reconstruct prev[i] and next[i]).
- Next, CAS all prev[i] to the new node. If a CAS fails, reconstruct prev[i].

Lock-Free Skiplists (Deletion)

- Pointers are atomically *markable*.
    - In C / C++, steal a bit from the pointer.
    - In Java, use AtomicMarkableReference.
- Construct the prev[i] and next[i] sets.
- For each pointer in our target node, mark them as deleted.

Optimized `decreaseKey` for Lock-Free Skiplists

- Currently, to implement a `decreaseKey` for Skiplists, we need to do two $\mathcal{O}(\log n)$ operations (one insert and one delete.)
- We can optimize this in some cases by removing redundant work: reuse the `prev[i]` set as a starting point for the next insertion.
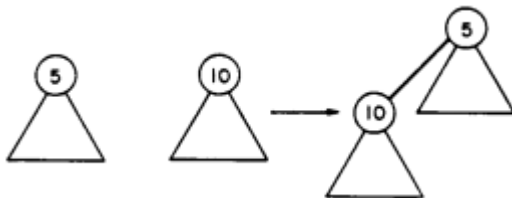
Pairing Heaps

- Pairing Heaps were introduced by Fredman and Tarjan[1] as a simplification of Fibonacci heaps.
  - Fibonacci heaps have a $\mathcal{O}(1)$ `decreaseKey` operator, but require complicated rebalancing.
- `decreaseKey` runs in $2^{\mathcal{O}(\sqrt{\log \log n})}$ time.
  - In practice, constant time. E.g., for a graph with a billion edges, $2^{\sqrt{\log \log 10^9}} = 3.34$.
- Pairing heaps have better performance than Fibonacci heaps on reasonably-sized graphs.[2]

---

[1] "The Pairing heap: a new form of self-adjusting heap," Fredman, Sedgewick, Sleator, and Tarjan, Algorithmica (1986).

[2] "On the Efficiency of Pairing Heaps and Related Data Structures," Michael Fredman, Journal of the ACM, 1999.
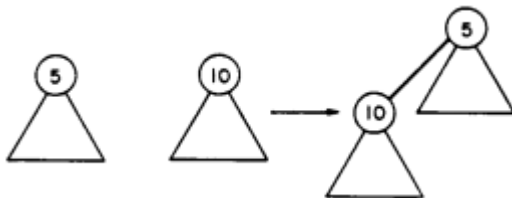
Pairing Heaps: Melding Two Heaps

- To meld two heaps, compare their roots and add the root as a child of the smaller root.
- meld is the core operator for Pairing heaps.

Pairing Heaps: `insert`

- To insert a new element into a pairing heap, construct a trivial heap containing only that element and `meld` this to the root.
- (The root of the heap will change if the new element is lower.)

Pairing Heaps: `deleteMin`

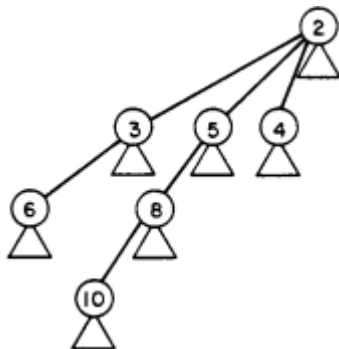- Pairing heaps have a two-step deleteMin operator.

Pairing Heaps: `deleteMin`

- After removing the root, `meld` each pair of heaps.

Pairing Heaps: `deleteMin`

- `meld` each resultant heap from right-to-left.

Parallelizing Dijkstra's Algorithm

- To simplify our Pairing heap implementation, we will only consider function interleavings that can occur when the heap is used for Dijkstra's algorithm.
- For Dijkstra's algorithm, we pop the minimum off the PQ (deleteMin) and call decreaseKey on the node's neighbors.
- If decreaseKey produces new heap minima , this will affect future calls to deleteMin
- For this reason, we assume that calls to decreaseKey and deleteMin never intersect.

Lock-Free Pairing Heaps: Overview

- Modify the list of sub-heaps to be a lock-free list of subheaps.
- We have identified an easy linearization point for `insert` and `deleteMin`: a node is considered in-place when the root is updated (CAS'd into place.)
  - But, what if the new node is inserted as a subheap of the current root?
  - We create a copy of the root to avoid this issue.

Lock-Free Pairing Heaps: `insert` an element

- Recall the insertion procedure: create a trivial heap containing only that element and `meld` it to the root.
- Lock-free `insert`:
  - Create a copy of the root, but the root and its copy share the same list of subheaps (this avoids an $\mathcal{O}(n)$ copy.)
  - Meld the new heap with the copy of the root.
  - Try to CAS the new root as the root of the heap.
  - If CAS succeeds, we're done.
  - If CAS fails, delete the subheap list changes and retry.

Lock-Free Pairing Heaps: `decreaseKey`

- Main observation: If Dijkstra's algorithm is run on a graph with exactly 0 or 1 edge between each node, calls to `decreaseKey` will be unique (per iteration.)
- **Case 1**: If the parent is still smaller than the node, just decrease the value.
    - We know that this node won't replace the root, since the parent is larger. An interleaved call to change the parent can only decrease its value.
- **Case 2**: If we're trying to decrease the value of the root, copy the root (as before), decrease its value, and try to CAS the new root into place. After a CAS fails, check if we are still trying to decrease the root and repeat if so.
- **Case 3**: The node is now smaller than its parent, violating the heap invariant. Immediately remove the node from the parent's list of subheaps. Next, follow the same procedure as `insert`.

Empirical Considerations

- Skiplists are extremely concurrent, but require two $\mathcal{O}(\log n)$ operations to implement decreaseKey. Also, operations require many pointer updates.

- Pairing heaps have a near-constant decreaseKey implementation and fewer pointer updates per call, but every operation is contending for the same location in memory (every thread contends for the root.)

Experimentation

- Run Dijkstra's algorithm on various large graphs (both randomly generated and real-world, e.g, from Harvard's Human Interactome Database.)
- Compare the time Dijkstra's algorithm takes for a Skiplist-backed implementation vs. a lock-free Pairing heap.
- Run experiments with varying numbers of threads.