

Facilitating Large-Scale Graph Searches with Lock-Free Pairing Heaps

Jeremy Mayeres, Charles Newton, Peter Tonner

12 March 2012

Project Overview

- We are constructing a lock-free version of a Pairing heap (a self-balancing heap.)
- Pairing heaps have an efficient decreaseKey implementation (near-constant performance) that allow you to decrease the value in a heap without reinserting it.
- This improves the asymptotic performance of certain algorithms (e.g., Dijkstra's algorithm.)
- We are comparing our heap against Skipqueues (a priority queue backed with a Skiplist.)

Dijkstra's Algorithm


Single-Source Shortest Path Problem

In a weighted graph $G = (V, E)$, find the shortest path from a target vertex $v \in V$ to all other vertices in the graph.

- Push every node in a graph into a priority queue (PQ.) In the PQ, each node is weighted by current information we have about their distance.
- Initially, all nodes have a distance of ∞ , except the target node, which has a distance of 0.
- Dynamic programming approach: Inductively build up our shortest routes.
 - Pop off the node on the PQ with the smallest distance. (The shortest path from the source to this node is finished.)
 - Update the weights in the PQ with the distances emanating from the popped node.


Skiplists

- Skiplists are constructed from a hierarchy of linked lists with the *Skiplist property*: the set of elements contained in level i is a subset of all the levels below it.
- Links allow you to do a binary search and “jump” around a list.
- The height of an element is randomly sampled from a power-law distribution: the probability of a node having a height of $i \geq 0$ is 2^{-i} .
- Time complexity of operations are probabilistically the same as for a binary search tree, but $\mathcal{O}(n)$ in the worst case.




Skiplists (Search)

- For each level, move right until you run into a node greater than your target. Then, from this point, move right on the next lowest level, and repeat.
- Likely runs in $\mathcal{O}(\log n)$ time.
- Head / tail nodes have values of $-\infty$ and ∞ , respectively.





Next and Previous Windows

- A useful abstraction to make is the set of all pointers related to a given node. (Aggregate all the pointers related to one node.)
- $\text{prev}[i] \rightarrow$ the node in level i pointing to the target node.
- $\text{next}[i] \rightarrow$ the node the target node points to at level i .
- Use the search process to construct these sets.



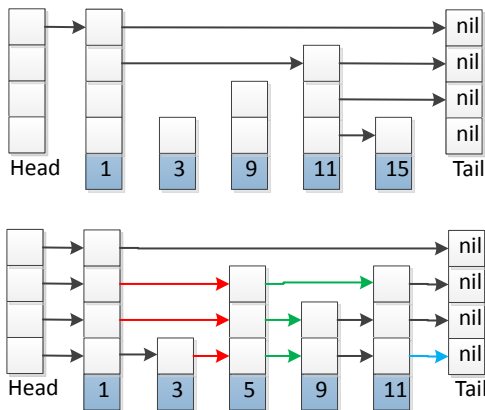
Skiplists (Insertion)

- Insertion: Find the location the target should be at. Set all of the node's next pointers to $\text{next}[i]$ and set the next pointers at each $\text{prev}[i]$ to the new node.



Skiplists (Deletion)

- Deletion: Set all of the next pointers in $\text{prev}[i]$ equal to $\text{next}[i]$.



Lock-free Skiplists

- Instead of linked lists at each level, we use lock-free linked lists.
- We lose the Skiplist property. (Levels aren't necessarily subsets of each other.) In particular, this means we need to always verify a node is in the lowest level.

Lock-free Skiplists

- Instead of linked lists at each level, we use lock-free linked lists.
- Linearization point: membership is defined at the lowest level of the Skiplist.
- We lose the Skiplist property. (Levels aren't necessarily subsets of each other.) In particular, this means we need to always verify a node is in the lowest level.

Lock-Free Skiplists (Insertion)

- Construct the $\text{prev}[i]$ and $\text{next}[i]$ sets.
- Set all of the node's next pointers to $\text{next}[i]$.
- If we can CAS the node into the bottom level, continue. Otherwise, something changed, and restart (reconstruct $\text{prev}[i]$ and $\text{next}[i]$).
- Next, CAS all $\text{prev}[i]$ to the new node. If a CAS fails, reconstruct $\text{prev}[i]$.

Lock-Free Skiplists (Deletion)

- Pointers are atomically *markable*.
 - In C / C++, steal a bit from the pointer.
 - In Java, use `AtomicMarkableReference`.
- Construct the `prev[i]` and `next[i]` sets.
- For each pointer in our target node, mark them as deleted.

Optimized decreaseKey for Lock-Free Skiplists

- Currently, to implement a decreaseKey for Skiplists, we need to do two $\mathcal{O}(\log n)$ operations (one insert and one delete.)
- We can optimize this in some cases by removing redundant work: reuse the `prev[i]` set as a starting point for the next insertion.

Pairing Heaps

- Pairing Heaps were introduced by Fredman and Tarjan¹ as a simplification of Fibonacci heaps.
 - Fibonacci heaps have a $\mathcal{O}(1)$ decreaseKey operator, but require complicated rebalancing.
- decreaseKey runs in $2^{\mathcal{O}(\sqrt{\log \log n})}$ time.
 - In practice, constant time. E.g., for a graph with a billion edges, $2^{\sqrt{\log \log 10^9}} = 3.34$.
- Pairing heaps have better performance than Fibonacci heaps on reasonably-sized graphs.²

¹“The Pairing heap: a new form of self-adjusting heap,” Fredman, Sedgewick, Sleator, and Tarjan, Algorithmica (1986).


²“On the Efficiency of Pairing Heaps and Related Data Structures,” Michael Fredman, Journal of the ACM, 1999.

Pairing Heaps

- Pairing Heaps were introduced by Fredman and Tarjan.
- Each node is a sub-heap
- The parent of any node has key smaller than the child node
- Each node has a left pointer to its first child, and a right pointer to its sibling
- Like decreaseKey, other operations run quickly as well
 - meld: $\mathcal{O}(1)$
 - findMin: $\mathcal{O}(1)$
 - insert: $\mathcal{O}(1)$
- delete and deleteMin run in $\mathcal{O}(\log n)$ amortized time

Pairing Heaps: Melding Two Heaps

- To meld two heaps, compare their roots and add the root as a child of the smaller root.
- `meld` is the core operator for Pairing heaps.
- Example: Melding two min-heaps together (root 3 and root 5)



Pairing Heaps: insert

- To insert a new element into a pairing heap, construct a trivial heap containing only that element and `meld` this to the root.
- (The root of the heap will change if the new element is lower.)

Pairing Heaps: decreaseKey

- Change the value of a node in the pair heap (e.g. updating the minimum distance to a node in Dijkstra's algorithm)
- If new value is still greater than parent node value, change in place
- Alternatively, remove node from heap, insert node with new value into heap

Pairing Heaps: deleteMin

- deleteMin removes the root node of a Pair Heap (e.g. pop off node with shortest distance in Dijkstra's algorithm)
- The root of the heap is first removed
- The heap property must then be restored by finding a new root
- Build pairs of heaps from the children of the root
- meld each resultant heap from right-to-left.

Pairing Heaps: deleteMin Pairing Operation


- The root of the tree has n sub-trees H_1, H_2, \dots, H_n , each with their own individual roots.
- For each tuple in the set $\{(H_1, H_2), (H_3, H_4), \dots, (H_{n-1}, H_n)\}$, we meld the two heaps together to form $\{H'_1, \dots, H'_m\}$.
- Perform a right-iteration (similar to foldr) of meld over these $m = \frac{n}{2}$ (or $m = \frac{n-1}{2} + 1$) heaps to construct a newly rebalanced pairing heap
- H' :

$$H' = \text{meld}(H'_1, \text{meld}(H'_2, \dots \text{meld}(H'_{m-1}, H'_m) \dots))$$


Pairing Heaps: deleteMin Example

- Example of the deleteMin operation


(a)



(b)




(c)



(d)



(e)




Parallelizing Dijkstra's Algorithm

- To simplify our Pairing heap implementation, we will only consider function interleavings that can occur when the heap is used for Dijkstra's algorithm.
- For Dijkstra's algorithm, we pop the minimum off the PQ (`deleteMin`) and call `decreaseKey` on the node's neighbors.
- If `decreaseKey` produces new heap minima, this will affect future calls to `deleteMin`
- For this reason, we assume that calls to `decreaseKey` and `deleteMin` never intersect.

Lock-Free Pairing Heaps: Overview

- Modify the list of sub-heaps to be a lock-free list of subheaps.
- We have identified an easy linearization point for `insert` and `deleteMin`: a node is considered in-place when the root is updated (CAS'd into place.)
 - But, what if the new node is inserted as a subheap of the current root?
 - We create a copy of the root to avoid this issue.
- Each node reference is updateable through a CAS operation
- In Java: `AtomicReference`
- Operations requiring linking nodes together use CAS loops
- Operations requiring node removal additionally use time stamps
- Build more complex lock-free operations from simpler ones

Lock-Free Pairing Heaps



Lock-Free Pairing Heaps: insert an element

- Recall the insertion procedure: create a trivial heap containing only that element and meld it to the root.
- Lock-free insert:
 - Create a copy of the root, but the root and its copy share the same list of subheaps (this avoids an $\mathcal{O}(n)$ copy.)
 - Meld the new heap with the copy of the root.
 - Try to CAS the new root as the root of the heap.
 - If CAS succeeds, we're done.
 - If CAS fails, delete the subheap list changes and retry.

Lock-Free Pairing Heaps: decreaseKey

- Main observation: If Dijkstra's algorithm is run on a graph with exactly 0 or 1 edge between each node, calls to decreaseKey will be unique (per iteration.)
- **Case 1:** If the parent is still smaller than the node, just decrease the value.
 - We know that this node won't replace the root, since the parent is larger. An interleaved call to change the parent can only decrease its value.
- **Case 2:** If we're trying to decrease the value of the root, copy the root (as before), decrease its value, and try to CAS the new root into place. After a CAS fails, check if we are still trying to decrease the root and repeat if so.
- **Case 3:** The node is now smaller than its parent, violating the heap invariant. Immediately remove the node from the parent's list of subheaps. Next, follow the same procedure as insert.

Lock-Free deleteMin

- Requires the removal of root node and restructuring of heap architecture
- Root reference and timestamp will be used to update the heap

Practical Considerations

- Skiplists are extremely concurrent, but require two $\mathcal{O}(\log n)$ operations to implement `decreaseKey`. Also, operations require many pointer updates.
- Pairing heaps have a near-constant `decreaseKey` implementation and fewer pointer updates per call, but every operation is contending for the same location in memory (every thread contends for the root.)

Experimentation

- Run Dijkstra's algorithm on various large graphs (both randomly generated and real-world, e.g, from Harvard's Human Interactome Database.)
- Compare the time Dijkstra's algorithm takes for a Skiplist-backed implementation vs. a lock-free Pairing heap.
- Run experiments with varying numbers of threads.